

AMATH 483 / 583 - Final

Due 5pm Friday June 9

May 31, 2023

Take Home Final (150 points, 20 extra credit points (EC))

1. (+20) **Fourier transforms.** Evaluate the Fourier transform of the following functions by hand. Use the definitions I provided (includes $\frac{1}{\sqrt{2\pi}}$, this is common in physics but also now the default used in WolframAlpha - a powerful math AI tool) as well as the definition for Dirac delta I used if needed.

(a) $f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2\sigma^2}(x-\mu)^2}$

(b) $f(t) = \sin(\omega_0 t)$

(c) $f(x) = e^{-a|x|}$ and $a > 0$

(d) (distribution) $f(t) = \delta(t)$

2. (+10) **Correlation.** By definition, *correlation* is $p \odot q = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} p^*(\tau)q(t+\tau)d\tau$, and measures how similar one signal or data function is to another. Let $p(\tau) = \langle p \rangle + \delta_p(\tau)$ and $q(\tau) = \langle q \rangle + \delta_q(\tau)$, where $\langle \rangle$ and $\delta(\cdot)$ denote the mean values and fluctuation functions (deviations about the mean). Two functions are defined to be *uncorrelated* when $p \odot q = \langle p \rangle \langle q \rangle$. Evaluate $p \odot q$ of the following functions:

$$p(t) = \begin{cases} 0 & t < 0 \\ 1 & 0 < t < 1 \\ 0 & t > 1 \end{cases}, \quad q(t) = \begin{cases} 0 & t < 0 \\ 1-t & 0 < t < 1 \\ 0 & t > 1 \end{cases}$$

3. (+5EC) **Autocorrelation.** Aside, periodic functions exhibit pronounced *autocorrelations* as shifting such functions by their period puts the function directly on itself. Alternatively, random functions or noise is characterized as being uncorrelated. Evaluate the autocorrelation $p \odot p$ of the following function:

$$p(t) = \begin{cases} 0 & t < 0 \\ 1 & 0 < t < 1 \\ 0 & t > 1 \end{cases}$$

4. (+20) **Fourier transform diffusion equation solve.** Consider the diffusion equation $\frac{\partial T}{\partial t} = \kappa \frac{\partial^2 T}{\partial x^2}$ where $T(x, t)$ describes the temperature profile of a long metal rod.

(a) Assume you know $T(x, 0)$ and define the Fourier transform of $T(x, t)$ to be $\tau(k, t)$. Transform the original equation and initial conditions into k -space. Solve the resulting equation. Inverse transform the result to obtain the solution in terms of the original variables.

(b) Find the temperature in the rod given initial conditions $\kappa = 10^3 \frac{m^2}{s}$ and

$$T(x, 0) = \begin{cases} 0 & |x| > 1m \\ 100^\circ\text{C} & |x| \leq 1m \end{cases}.$$

5. (+20) **Compare OpenBLAS to CUBLAS on HYAK.** Measure and plot the performance of double precision matrix multiply ($\alpha AB + \beta C \rightarrow C$) for square matrices of dimension $n = 16$ to $n = 8192$, stride $n* = 2$ for both the OpenBLAS and CUDA BLAS (CUBLAS) implementations on HYAK. Let each n be measured $ntrial$ times and plot the average performance for each case versus n , $ntrial \geq 3$. Submit your performance plot and C++ test code. Your plot will have 'mflops' on the y-axis and the dimension of the matrices on the x-axis.

6. (+10EC) **Matrix transpose.** Write C++ functions given the two APIs that compute A^T , the matrix transpose. Test the correctness of both functions. Put the functions in file **transpose.hpp** which I will include in my test code for grading. Submit file **transpose.hpp**. Note the threaded function will create and join the threads internally.

```
void sequentialTranspose(std::vector<int> &matrix, int rows, int cols);
void threadedTranspose(std::vector<int> &matrix, int rows, int cols, int nthreads);
```

7. (+20) **Memory access time.** On a computer of your choice, write C++ functions for the given APIs that perform row and column swap operations in *memory* on a type double matrix stored in column major index order using a single vector container for the data. Test the swap capabilities on randomly selected index pairs using the function I provide here. Put your functions (not 'getRandomIndices') in file **mem_swaps.hpp**, no need for header guards -just the code for your functions. Conduct a performance test for square matrix dimensions 16, 32, 64, 128, ... 16384, measuring the time required to conduct row and column swaps separately. Let each operation be measured *ntrial* times and plot the average time versus matrix dimension, $ntrial \geq 3$. Make a single plot of your *row* and *column* swap timing measurements with time on the y-axis ($\log_{10}(time)$) and the problem dimension on the x-axis. Submit file **mem_swaps.hpp** and plot.

```
void swapRows(std::vector<double> &matrix, int nRows, int nCols, int i, int j);
void swapCols(std::vector<double> &matrix, int nRows, int nCols, int i, int j);
```

```
#include <utility> // For std::pair
std::pair<int, int> getRandomIndices(int n)
{
    int i = std::rand() % n;
    int j = std::rand() % (n - 1);
    if (j >= i)
    {
        j++;
    }
    return std::make_pair(i, j);
}

// ... from inside main()
//std::pair<int, int> rowIndices = getRandomIndices(M);
//int i = rowIndices.first;
//int j = rowIndices.second;
//std::pair<int, int> colIndices = getRandomIndices(N);
// ...
```

8. (+20) **File access time.** On the same computer as problem 1, write C++ functions for the given APIs that perform row and column swap operations on a type double matrix stored in column major index order in a *FILE*. Use a randomly selected pair of indices to test the swapping capabilities. Put the functions you write in file **file_swaps.hpp**. Conduct a performance test for square matrix dimensions 16, 32, 64, 128, ... 16384, measuring the time required to conduct *file-based* row and column swaps separately. Let each operation be measured *ntrial* times and plot the average time versus matrix dimension, $ntrial \geq 3$. Make a single plot of your *file-based row* and *column* swap timing measurements with time on the y-axis ($\log_{10}(time)$) and the problem dimension on the x-axis. Submit your header file **file_swaps.hpp** and plot.

```
void swapRowsInFile(std::fstream &file, int nRows, int nCols, int i, int j);
void swapColsInFile(std::fstream &file, int nRows, int nCols, int i, int j);
```

```
// snippet
#include <iostream>
#include <fstream>
#include <vector>
#include <utility>
#include <algorithm>
#include <cstdlib>
```

```

#include <ctime>
#include <cstdio>
#include <chrono>
#include "file_swaps.hpp"

int main(int argc, char *argv[])
{
    // Generate the matrix
    std::vector<double> matrix(numRows * numCols);
    // init matrix elements in column major order
    // write the matrix to a file
    std::fstream file(filename, std::ios::out | std::ios::binary);
    file.write(reinterpret_cast<char *>(&matrix[0]), numRows * numCols * sizeof(double));
    file.close();
    // Open the file in read-write mode for swapping
    std::fstream fileToSwap(filename, std::ios::in | std::ios::out | std::ios::binary);
    // Get random indices i and j for row swapping
    // Measure the time required for row swapping using file I/O
    auto startTime = std::chrono::high_resolution_clock::now();
    // Swap rows i and j in the file version of the matrix
    swapRowsInFile(fileToSwap, numRows, numCols, i, j);
    auto endTime = std::chrono::high_resolution_clock::now();
    std::chrono::duration<double> duration = endTime - startTime;
    // Close the file after swapping
    fileToSwap.close();
    //...
    // after each problem size delete the test file
    std::remove(filename.c_str());
    // ...
}

```

9. (+20) **CPU-GPU data copy speed on HYAK.** Write a C++ code to measure the data copy performance between the host CPU and GPU, and between the GPU and the host CPU. Copy 8 bytes to 256MB increasing in multiples of 2. You will plot the bandwidth for both directions: (bytes per second) on the y-axis, and the buffer size in bytes on the x-axis. Submit your plot and test code.
10. (+20) **Compare FFTW to CUFFT on HYAK.** Measure and plot the performance of calculating the derivative of a 3D double complex plane wave defined on cubic lattices of dimension n^3 from 16^3 to $n = 512^3$, stride $n* = 2$ for both the FFTW and CUDA FFT (CUFFT) implementations on HYAK. Let each n be measured $ntrial$ times and plot the average performance for each case versus n , $ntrial \geq 3$. Submit your performance plot and C++ test code. Your plot will have 'mflops' on the y-axis and the dimension of the cubic lattices (n) on the x-axis. You will need to estimate the operation count of computing the derivative using FFT on a lattice.
11. (+5EC) **Root finding.** Write a C++ function that implements a Newton or bisection iteration to estimate the *real* roots to a polynomial equation. Your code should accept the degree of the polynomial, the coefficients, and the domain from the command line (as below). Submit your test code. I will run it against a couple test polynomials of degree 3 and 4.

```

Enter the degree of the polynomial: 3
Enter coefficient 3: *
Enter coefficient 2: *
Enter coefficient 1: *
Enter coefficient 0: *
Enter the start of the domain: 0
Enter the end of the domain: 10
Roots found:
-0.*****
1.**
3.*****

```